

DataONE: A Data Federation with Provenance Support

Yang Cao¹, Christopher Jones², Víctor Cuevas-Vicentín³, Steve Aulenbach⁴,
Matthew B. Jones², Bertram Ludäscher¹, Timothy McPhillips¹, Paolo Missier⁵,
Christopher Schwalm⁶, Peter Slaughter², Dave Viegals², Lauren Walker²,
Yaxing Wei⁷

¹ Library and Information Science, University of Illinois, Urbana-Champaign, IL

² National Center for Ecological Analysis and Synthesis, UCSB, CA

³ Universidad Popular Autónoma del Estado de Puebla, Mexico

⁴ UCAR and U.S. Global Change Research Program

⁵ School of Computing Science, Newcastle University, UK

⁶ Woods Hole Research Center, Falmouth, MA

⁷ Environmental Sciences Division, Oak Ridge National Laboratory, TN

Abstract. DataONE is a federated data network focusing on earth and environmental science data. We demonstrate new provenance capabilities in the DataONE toolkit to facilitate reproducible research. A user “Alice”, can annotate a (Matlab, R, etc.) script using the YesWorkflow (YW) tool to describe the underlying workflow or *prospective* provenance. After Alice has run the script, the result files, script, prospective provenance, and *retrospective* provenance, represented in the ProvONE provenance model, are bundled into an OAI-ORE compliant data package and uploaded to the DataONE network. A second user (“Bob”) discovers Alice’s package and uses her data in his own analysis. We show that Bob’s results, once published through DataONE, link back to Alice’s outputs via unique identifiers. Thus, a third user (“Charlie”) who browses DataONE discovers the full provenance of Bob’s results, all the way back to Alice’s original contributions. DataONE provenance systems enable reproducible research and facilitate proper attribution of scientific results transitively across generations of derived data products.

1 Introduction

Scientific workflow *provenance* is valuable in computational science. Provenance can help scientists to understand their own work and share their work with others while maintaining attribution. We refer to two types of provenance: *prospective* and *retrospective* provenance.

DataONE (Data Observation Network for Earth) is a federated data network and a sustainable cyberinfrastructure for open, persistent, robust, and secure access to well-described and easily discovered Earth observational data [3]. The primary goals of DataONE are: discovery and access, data integration and synthesis, education and training, building community, and data sharing. The DataONE infrastructure consists of three principal components:

(1) **Member Nodes** are existing or new data repositories that support the DataONE Member Node API. Tier-1 Member Nodes (MN) implement anonymous, read-only access to data and metadata and enable discovery of all objects available on the MN, with

low level descriptions of each object, object retrieval by identifier, and activity reporting. MNs also provide science metadata and relationships between metadata and data using resource maps [12] to facilitate discovery. MNs can also implement higher tiers of the DataONE API, including authenticated and restricted data access (Tier-2), write access to storage resources (Tier-3), and replication of data from other nodes (Tier-4).

(2) Coordinating Nodes (CN) serve coordination and discovery needs of the network. Services of a CN include network-wide indexing of data objects, coordination of data replication between MNs, mirrored content of *science metadata* (detailed descriptions of science data objects and collections) and *system metadata* (low level metadata, e.g., data type, size, ownership, location) present at MNs.

(3) Investigator Toolkit. The Investigator Toolkit (ITK) contains user tools that enable interaction with DataONE infrastructure through a REST service API exposed by the CNs and MNs. Python and Java libraries are available for application developers, and DataONE support is present in scientific analysis tools such as Matlab and R.

The DataONE infrastructure was released for public use in 2012 and supports identifier resolution, content search and retrieval, the federated identity management infrastructure, and the replication service.

DataONE Search is a web-based application allowing users to seamlessly and efficiently discover publicly accessible data packages within the DataONE federated network of Member Node repositories. It allows users to search by map, dates, keywords and other facets and to refine the results using further parameters. Users sign in to DataONE Search using ORCID credentials, Google accounts, or institutional accounts.

The software has been designed to facilitate rapid iteration and deployment of new features and to take full advantage of future capabilities offered by upcoming versions of the core cyberinfrastructure. Notable among these is the *provenance information* within search and discovery. In the next section, we will demonstrate new DataONE provenance tools and the visualization of provenance with DataONE Search.

2 Demonstration Description

DataONE has developed a *Run Manager* API for capturing provenance from R [11] and Matlab [5] script executions, and contributed to the development of YesWorkflow [8] to provide an improved understanding of a script in terms of its user-declared workflow structure and system-recorded file I/O. The automatic capture of the latter is transparent to scientists and provides medium-grained⁸ provenance by focusing on file I/O events.

We will use an example to introduce the core provenance and search features of DataONE. Suppose that Alice is an earth scientist whose work involves generating global grass fraction maps from water, air temperature, and other input data. Bob is a scientist who runs a grassland water use efficiency analysis built upon Alice's output. Charlie is a user of DataONE.

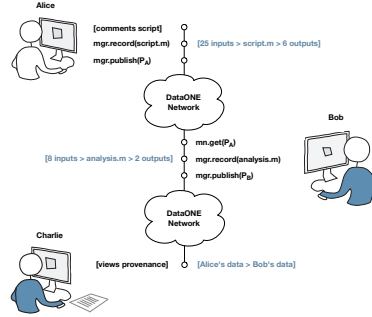


Fig. 1: Run Manager Demonstration: (1) Alice runs `script.m` with the DataONE Run Manager to create data package P_A , which she publishes to the DataONE network; (2) Bob later finds and downloads Alice’s data, uses it in his `analysis.m`, creating and then publishing package P_B ; (3) Charlie searches DataONE, finds Bob’s P_B , and recognizes its dependence on Alice’s P_A .

2.1 User Alice’s view: Local Run Manager View

In Figure 1, Alice has developed a script for producing Carbon3/Carbon4 soil maps. She uses the YesWorkflow (YW) tool to mark-up the script and expose the underlying workflow view (i.e., prospective provenance) that is inherent in her soil mapping code. Alice then runs her script using the *Run Manager* (the DataONE provenance capture tool) in Matlab. After a couple of runs, Alice is happy with the results and publishes them to the DataONE network. To do so, she uses the tool to bundle up the data results along with the runtime (retrospective) provenance captured by the DataONE Matlab provenance recorder, the script itself and its YW-generated workflow view.

The DataONE Matlab tool facilitates her work by generating the DataONE-compliant data package in OAI-ORE format, including the provenance information in ProvONE format as part of the package. ProvONE [13] is an extension of the W3C PROV-O [14] standard for representing provenance, and includes specializations for representing both retrospective provenance about the runtime execution and prospective provenance about the structure and flow of the analytical script or workflow.

When using DataONE Search with the keyword “grass”, Bob finds Alice’s data package [2]. Bob uses outputs from Alice for his own study, which he also shares via DataONE. Since he used the unique identifiers from Alice’s outputs, the provenance traces produced by Alice’s and Bob’s runs are naturally connected through the nodes that carry those data identifiers. Charlie discovers Bob’s data packages on DataONE and is able to navigate back to the data that Bob used, i.e., Alice’s data package depicted in Figure 3 [6,10].

2.2 User Bob’s view: Local Run Manager View

When Bob searches for data in DataONE and finds Alice’s data and code, he uses the Matlab Toolbox to fetch the data in a way that is identifier-aware, rather than

⁸ unlike the fine-grained provenance captured, e.g., by noWorkflow [7]

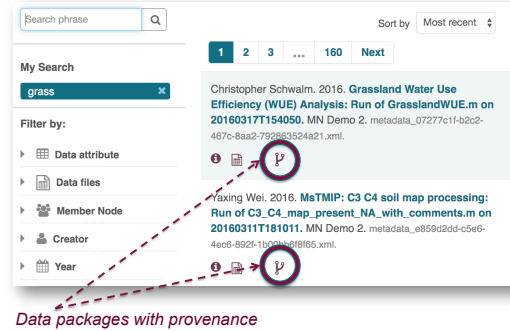


Fig. 2: DataONE Search result for keyword “grass” (DataONE Search demonstration site). The “fork” icon indicates data packages having embedded provenance.

having to manually maintain the provenance information. For example, he calls the `MemberNode.get(session, identifier)` function in his analysis script because this function has provenance tracking capabilities. Then, the Toolbox downloads the data and uses the same identifier to store the data locally. Additionally, a *prov:used* statement that Bob is using Alice’s data is recorded.

When Bob then calls `publish()` to upload his scripts and data, data in his package that are not present in the DataONE network are uploaded. Their respective provenance statements are indexed on the *Coordinating Node*, and are available for use in the web provenance display.

2.3 User Charlie’s View: DataONE Search Website

When Charlie searches the keyword “grass” in the DataONE Search, two data packages are found and shown in Figure 2. One data package is created by “Alice” (Yaxing Wei in the real world) [2], the other is created by “Bob” (Christopher Schwalm) [1]. Both data packages show that provenance information is associated with each dataset (via the icon in the search record) and can be seen within the DataONE Search demo site [4].

Prospective and retrospective provenance information can be explored from the DataONE Search site. Provenance details for any input or output in the provenance graph can be viewed by clicking on the icon shown in Figure 3. DataONE Search provides human language descriptions of how data are used or generated via the script and models, and provides navigation to ancestors and descendants in the data derivation chain. Figure 3 shows two data packages produced by the execution of Alice and Bob’s workflow scripts, and the provenance lineage between these two data packages.

Alice’s data package [2] is shown on the top layer of Figure 3, created when Alice ran her script using the DataONE provenance capture tool, or Run Manager, in Matlab. Her script (`C3_C4_map_present_with_comments.m`) takes twenty-five input files and produces six outputs, which are shown on the left side and right side of Alice’s data package in Figure 3, respectively. The bottom three outputs in Alice’s data package are the NetCDF data files that represent three different world map grids of percentage of grass types (C3 grass fraction, C4 grass fraction, and total grass fraction). In addition, a

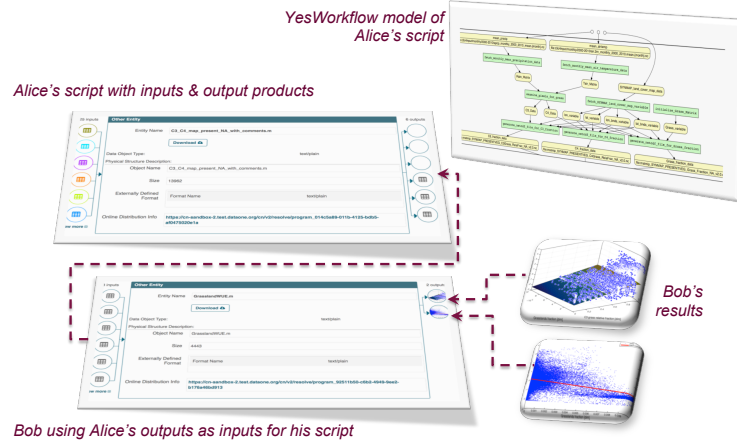


Fig. 3: Charlie's view on the DataONE demo site: (1) A YesWorkflow model for Alice's soil processing script; (2) Data lineage from Bob's results back through his script inputs to Alice's data package; (3) Two visualizations produced by Bob's water use efficiency analysis script.

model graph is displayed at the intermediate layer that was generated by the YesWorkflow tool [8]. Alice's used embedded YesWorkflow annotations to document her script. These annotations declare step by step how data are used and derived in the script.

Bob's data package [1] created by DataONE provenance tool (Run Manager) is displayed at the bottom layer. When Bob browses Alice's data package on the DataONE Search site, he decides to use three NetCDF output data produced by Alice's work in his Grassland Water Use Efficiency Analysis script (GrasslandWUE.m). From the provenance information associated with Bob's data package in Figure 3, we see that it takes eight inputs, and produces two visualizations. By viewing the details for each input, we can see that three of them are the outputs produced by Alice's data package.

In order to maintain the link between the outputs from Alice's data package to the inputs to Bob's package, Bob needs to use the same identifier in his script as Alice does. An alternative approach is that Bob creates a new identifier when he downloads Alice's data. However, the link between two data packages will be broken which has been discussed in [9].

3 Discussion and future work

This paper demonstrates a notable new feature (provenance capture and search) of DataONE. DataONE has released the DataONE Search for public use in November 2015, and the R and Matlab provenance tool to the public in 2016.

Maintaining the link between Alice's work and Bob's subsequent work is worth discussion. Currently, Bob must use certain functions provided by DataONE to ensure the provenance link to Alice's work is correctly maintained. There are other possible ways to achieve the same goal. For example, Bob might download Alice's data and

manually add the links to Alice's work back into his data package before sharing via DataONE. A *prov:was_derived_from* statement that local data copy for Bob is derived from Alice's data could easily add, especially for historical data for which provenance was not captured at the time data were generated. The new provenance tools described in this demonstration allow efficient capture of machine-parseable provenance information using analytical tools commonly used in the environmental sciences (e.g., Matlab), thereby significantly improving the replicability of research in the environmental sciences.

Future work of the current provenance tools include: (1) solve the broken link use case; (2) support more I/O functions; (3) handle complex scenarios such as multiple runs of a script and multiple users; (4) variable-level provenance capture warrants further investigation and efforts.

References

1. Christopher Schwalm's DataPackage, https://search-sandbox-2.test.dataone.org/#view/metadata_07277c1f-b2c2-467c-8aa2-792863524a21.xml
2. YaXing Wei's DataPackage, https://search-sandbox-2.test.dataone.org/#view/metadata_e859d2dd-c5e6-4ec6-892f-1b00bb6f8f65.xml
3. Data Observation Network for Earth (DataONE), <https://www.dataone.org>
4. DataONE Search Demo Site, <https://search-sandbox-2.test.dataone.org>
5. Jones C., Cao Y., Slaughter P., Jones M.: Matlab DataONE Toolbox. <https://github.com/DataONEorg/matlab-dataone> (2016)
6. Katz, D.S. & Smith, A.M.: Implementing Transitive Credit with JSON-LD. *Journal of Open Research Software*. 3(1), p.e7. (2015)
7. Murta, L., Braganholo, V., Chirigati, F., Koop, D., Freire, J.: noWorkflow: Capturing and analyzing provenance of scripts. In: *International Provenance and Annotation Workshop (IPAW)*. pp. 71–83. Cologne, Germany (2014)
8. McPhillips T., Song T., Kolisnik T., Aulenbach S., Belhajjame K., Bocinsky R.K., Cao Y., Cheney J., Chirigati F., Dey S., Freire J., Jones C., Hanken J., Kintigh K.W., Kohler T.A., Koop D., Macklin J. A., Missier P., Schildhauer M., Schwalm C., Wei Y., Bieda M., and Ludäscher B.: YesWorkflow: A User-Oriented, Language-Independent Tool for Recovering Workflow Information from Scripts. *Intl. Journal for Digital Curation*, 10, 298–313. (2015)
9. Missier, P., Ludäscher, B., Bowers, S., Anand, M. K., Altintas, I., Dey, S., Sarkar, A., Shrestha, B. and Goble, C.: Linking Multiple Workflow Provenance Traces for Interoperable Collaborative Science. In: *5th Workshop on Workflows in Support of Large-Scale Science*. (2010).
10. Missier, P.: Data Trajectories: Tracking Reuse of Published Data for Transitive Credit Attribution. In: *11th International Data Curation Conference*. (2016).
11. Slaughter P., Jones M., Jones C.: NCEAS recordr: Provenance tracking for R. <https://github.com/NCEAS/recordr> (2016)
12. Open Archives Initiative Object Reuse and Exchange, <https://www.openarchives.org/ore/>
13. ProvONE: A PROV Extension Data Model for Scientific Workflow Provenance, <https://purl.dataone.org/provone-v1-dev>
14. W3C PROV-O: The PROV Ontology, <https://www.w3.org/TR/prov-o/>